

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Random Forest Method

Gini Impurity Calculation

Gini impurity measures how mixed the classes are within a node. It is computed as:

$$\text{Gini impurity}(S) = 1 - \sum_{i=1}^k p_i^2$$

Where:

- S = collection of features
- p_i = proportion of samples belonging to class (i)
- k = total number of classes

A Gini value of 0 indicates a perfectly pure node (all samples belong to one class), while higher values indicate more class mixing.

Information Gain (Split Criterion)

Decision trees choose the split that produces the largest reduction in impurity.

Information gain is defined as:

$$\text{Gain}(S, A) = \text{Impurity}(S) - \sum_{v \in \text{Value}(A)} \frac{|S_v|}{|S|} \text{Impurity}(S_v)$$

Where:

- S is a collection of features
- A is a feature in S
- $\text{Value}(A)$ is the set of all possible values for attribute A
- S_v is the subset of S for which attribute A has value v .

The split with the highest information gain is selected.

Baseline Random Forest Model

This is run on the cleaned/normalized data as well as sklearn default parameters. This

gives us an idea of how this model may perform and potential hyperparameters to focus on during fine-tuning.

In [2]:

```
from sklearn.model_selection import train_test_split

# load data
df_norm = pd.read_csv("../..../data/processed/winequality-red-normalized.csv"

# split into features and target
X_norm = df_norm.drop(columns=['quality'])
y_norm = df_norm['quality']

X_train, X_test, y_train, y_test = train_test_split(
    X_norm, y_norm, test_size=0.2, random_state=42, stratify=y_norm
)
```

In [3]:

```
from sklearn.ensemble import RandomForestClassifier

# fit baseline random forest
rf_baseline = RandomForestClassifier(random_state=42)

rf_baseline.fit(X_train, y_train)
```

Out[3]: RandomForestClassifier(random_state=42)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [4]:

```
from sklearn.metrics import accuracy_score, f1_score

y_pred = rf_baseline.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred, average='macro'))
```

Accuracy: 0.684375
F1 Score: 0.40597237221330523

In [5]:

```
from sklearn.metrics import classification_report

report_dict = classification_report(y_test, y_pred, output_dict=True, zero_
df_report = pd.DataFrame(report_dict).transpose()
df_report
```

Out [5] :

	precision	recall	f1-score	support
3	0.000000	0.000000	0.000000	2.000000
4	0.000000	0.000000	0.000000	11.000000
5	0.725352	0.757353	0.741007	136.000000
6	0.643836	0.734375	0.686131	128.000000
7	0.724138	0.525000	0.608696	40.000000
8	0.500000	0.333333	0.400000	3.000000
accuracy	0.684375	0.684375	0.684375	0.684375
macro avg	0.432221	0.391677	0.405972	320.000000
weighted avg	0.661014	0.684375	0.669218	320.000000

The random forest baseline model performs well on the most common wine qualities (5 and 6), achieving F1 scores of above 0.66, meaning that it's doing a pretty good job at correctly identifying the class while balancing false positives and false negatives. It struggles more with the minority classes (3, 4, and 8) due to the class imbalance present. The overall accuracy of 68.4% is pretty strong as a baseline for this dataset, but further fine-tuning is needed to improve performance further. However, the overall macro F1 score of 0.40 indicates that the model is again, performing poorly across all classes, indicating issues with minority classes.

Improving the baseline model

Next, we will perform Grid Search Cross Validation on each of the three datasets (our original, PCA, and interaction term).

In [6] :

```
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, f1_score, classification_report
```

In [7] :

```
datasets = {
    "normalized": pd.read_csv("../data/processed/winequality-red-normalized.csv"),
    "pca": pd.read_csv("../data/processed/winequality-red-pca.csv"),
    "interactions": pd.read_csv("../data/processed/winequality-red-interactions.csv")
}
```

In [8] :

```
param_grid = {
    "n_estimators": [50, 100, 150, 200],
    "max_depth": [None, 15, 25, 50],
    "max_features": ['sqrt', 'log2', None],
    "class_weight": [None, "balanced"]
}
```

In [9]:

```
results_table = []

for name, df in datasets.items():
    print(f"\n{name} dataset")

    # split into X and y
    X = df.drop(columns=["quality"])
    y = df["quality"]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42, stratify=y
    )

    # grid search
    grid = GridSearchCV(
        RandomForestClassifier(random_state=42),
        param_grid,
        scoring="f1_macro",
        cv=5,
        n_jobs=-1,
        return_train_score=False
    )

    grid.fit(X_train, y_train)

    best_params = grid.best_params_
    best_cv_score = grid.best_score_
    cv_std = grid.cv_results_['std_test_score'][grid.best_index_]

    print("Best parameters:", best_params)
    print(f"CV Macro F1: {best_cv_score:.4f} (+/- {cv_std:.4f})")

    # train best
    best_model = RandomForestClassifier(**best_params, random_state=42)
    best_model.fit(X_train, y_train)

    # predict
    y_pred = best_model.predict(X_test)

    # metrics
    test_acc = accuracy_score(y_test, y_pred)
    test_macro_f1 = f1_score(y_test, y_pred, average="macro")

    print(f"Test Accuracy: {test_acc:.4f}")
    print(f"Test Macro F1: {test_macro_f1:.4f}\n")
    print("Classification Report:")
    print(classification_report(y_test, y_pred, zero_division=0))

    # normalized confusion matrix
    cm = confusion_matrix(y_test, y_pred, normalize="true")
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=sorte
    disp.plot(cmap="Blues", values_format=".2f")
    plt.title(f"Normalized Confusion Matrix - {name} dataset")
    plt.show()
```

```

# store results in table
results_table.append({
    "Dataset": name,
    "CV Macro F1 (mean)": round(best_cv_score, 4),
    "CV Macro F1 (std)": round(cv_std, 4),
    "Test Accuracy": round(test_acc, 4),
    "Test Macro F1": round(test_macro_f1, 4)
})

results_df = pd.DataFrame(results_table)

```

normalized dataset

Best parameters: {'class_weight': 'balanced', 'max_depth': None, 'max_features': None, 'n_estimators': 100}

CV Macro F1: 0.3523 (+/- 0.0400)

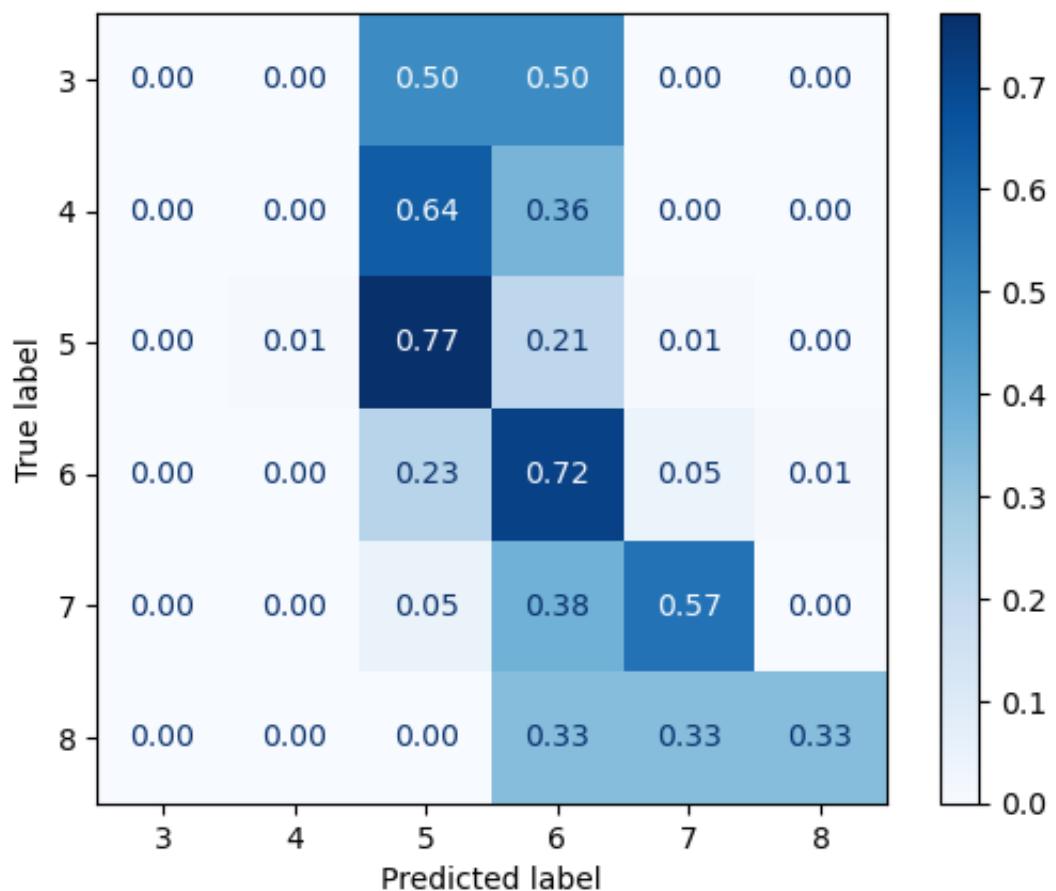
Test Accuracy: 0.6906

Test Macro F1: 0.4132

Classification Report:

	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	11
5	0.73	0.77	0.75	136
6	0.65	0.72	0.68	128
7	0.74	0.57	0.65	40
8	0.50	0.33	0.40	3
accuracy			0.69	320
macro avg	0.44	0.40	0.41	320
weighted avg	0.67	0.69	0.68	320

Normalized Confusion Matrix - normalized dataset

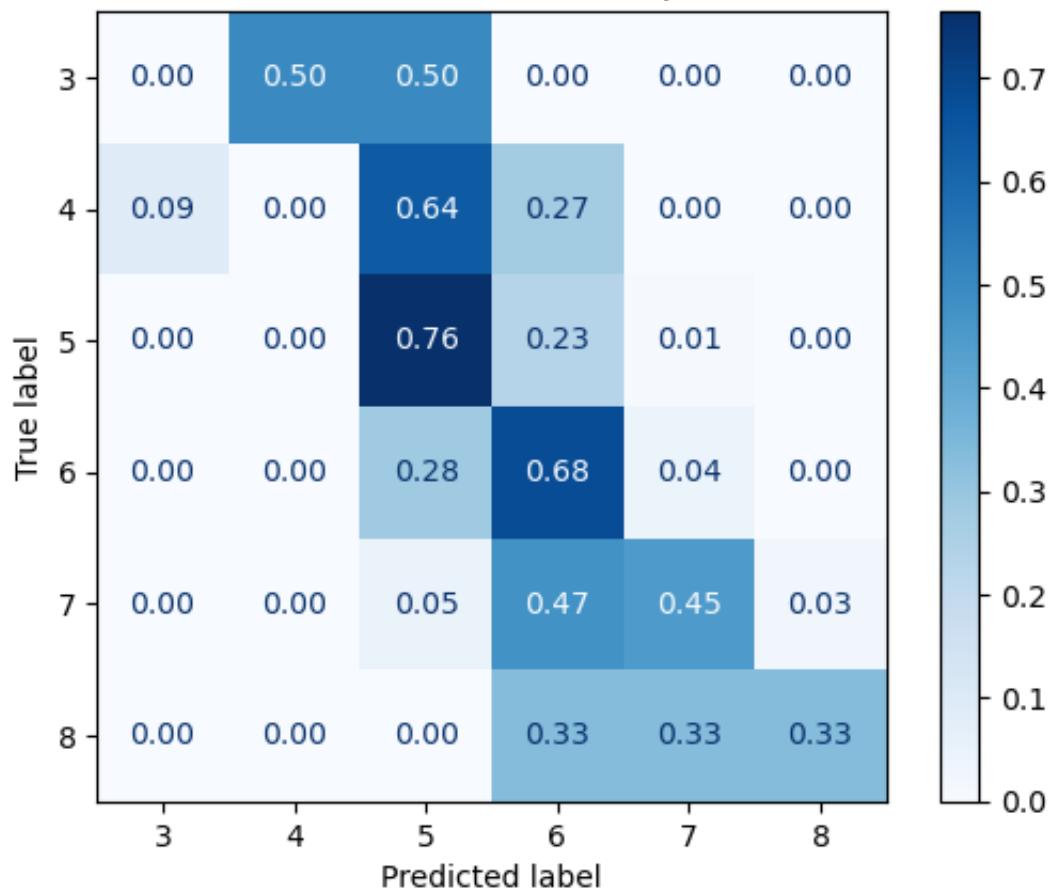


```
pca dataset
Best parameters: {'class_weight': None, 'max_depth': 15, 'max_features': 'sqrt', 'n_estimators': 50}
CV Macro F1: 0.3625 (+/- 0.0693)
Test Accuracy: 0.6562
Test Macro F1: 0.3880
```

Classification Report:

	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	11
5	0.69	0.76	0.73	136
6	0.62	0.68	0.65	128
7	0.72	0.45	0.55	40
8	0.50	0.33	0.40	3
accuracy			0.66	320
macro avg	0.42	0.37	0.39	320
weighted avg	0.64	0.66	0.64	320

Normalized Confusion Matrix - pca dataset



interactions dataset

Best parameters: {'class_weight': 'balanced', 'max_depth': 15, 'max_features': 'sqrt', 'n_estimators': 100}

CV Macro F1: 0.3477 (+/- 0.0491)

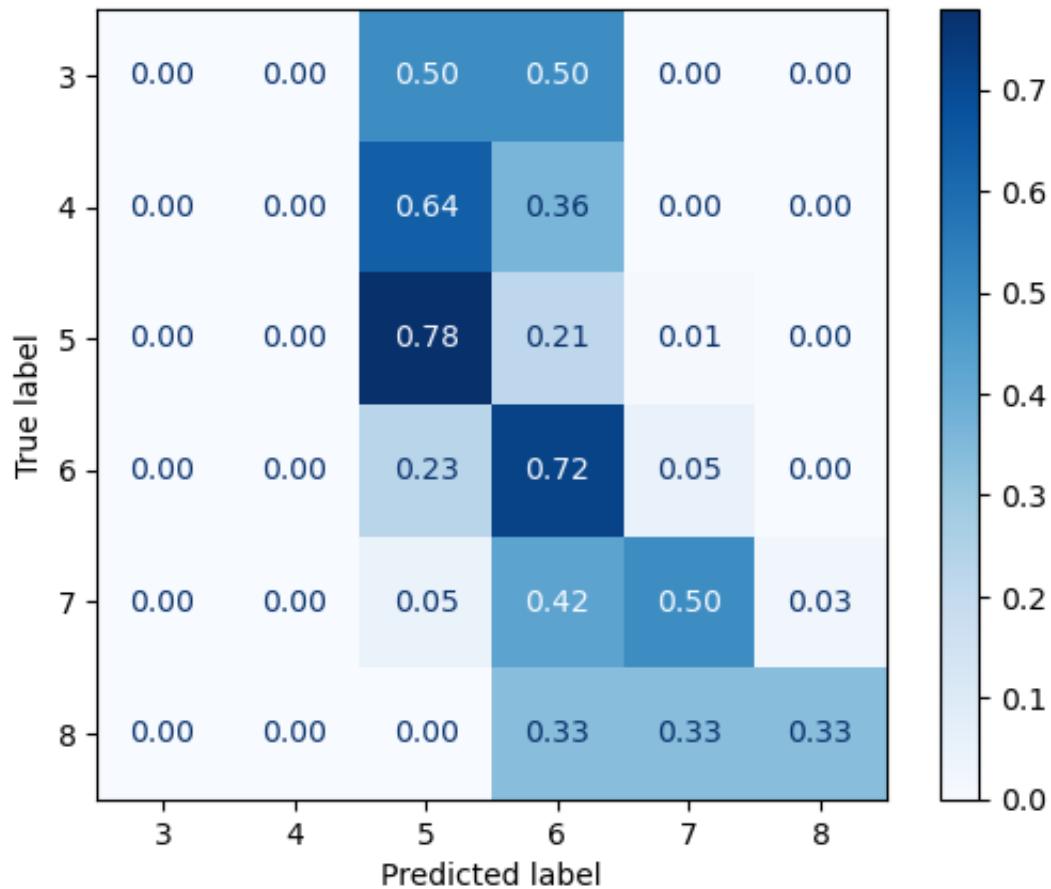
Test Accuracy: 0.6844

Test Macro F1: 0.4018

Classification Report:

	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	11
5	0.73	0.78	0.75	136
6	0.64	0.72	0.68	128
7	0.69	0.50	0.58	40
8	0.50	0.33	0.40	3
accuracy			0.68	320
macro avg	0.43	0.39	0.40	320
weighted avg	0.66	0.68	0.67	320

Normalized Confusion Matrix - interactions dataset



In [10]: `results_df`

	Dataset	CV Macro F1 (mean)	CV Macro F1 (std)	Test Accuracy	Test Macro F1
0	normalized	0.3523	0.0400	0.6906	0.4132
1	pca	0.3625	0.0693	0.6562	0.3880
2	interactions	0.3477	0.0491	0.6844	0.4018

Effects of hyperparameter class_weight

Here, I will test the efficacy of the `class_weight` hyperparameter on model performance on the normalized dataset.

In [11]: `X = df_norm.drop(columns=["quality"])`
`y = df_norm["quality"]`

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42, stratify=y  
)
```

```
# best hyperparameters from GridSearch besides class_weight
best_params_fixed = {
    "n_estimators": 100,
    "max_depth": None,
    "max_features": None,
    "random_state": 42
}

results = {}

for cw in [None, "balanced"]:
    print(f"\nUsing class_weight = {cw}")

    model = RandomForestClassifier(class_weight=cw, **best_params_fixed)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    macro_f1 = f1_score(y_test, y_pred, average="macro")

    print(f"Accuracy: {acc:.4f}")
    print(f"Macro F1: {macro_f1:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred, zero_division=0))
```

```
Using class_weight = None
```

```
Accuracy: 0.6719
```

```
Macro F1: 0.3888
```

```
Classification Report:
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	11
5	0.71	0.76	0.73	136
6	0.63	0.71	0.67	128
7	0.74	0.50	0.60	40
8	0.33	0.33	0.33	3
accuracy			0.67	320
macro avg	0.40	0.38	0.39	320
weighted avg	0.65	0.67	0.66	320

```
Using class_weight = balanced
```

```
Accuracy: 0.6906
```

```
Macro F1: 0.4132
```

```
Classification Report:
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	11
5	0.73	0.77	0.75	136
6	0.65	0.72	0.68	128
7	0.74	0.57	0.65	40
8	0.50	0.33	0.40	3
accuracy			0.69	320
macro avg	0.44	0.40	0.41	320
weighted avg	0.67	0.69	0.68	320

Without the class_weight parameter, the model achieves an accuracy of 0.67 and macro F1 of 0.389, indicating problems with performance in smaller classes. Introducing class weights improves macro F1 to 0.4132, indicating increase in performance on smaller classes (particularly quality = 8). This also did not sacrifice performance on the bigger classes, and improved accuracy overall.

```
In [ ]:
```